

Digital Normalization


Lecture 16
Oct 10, 2016

Announcements

Perfect Storm of data analysis

DIGITAL NORMALIZATION

Perfect Storm of data analysis – What to do???



Brown 2012 arXiv:1203.4802v2


DIGITAL NORMALIZATION

----- True sequence (unknown)


_____X_____

Reads
(randomly sequenced)

DIGITAL NORMALIZATION




DIGITAL NORMALIZATION



```
for read in dataset:  
    if estimated_coverage(read) < C:  
        accept(read)  
    else:  
        discard(read)
```


DIGITAL NORMALIZATION

True sequence (unknown)




Reads
(randomly sequenced)

```
for read in dataset:
    if estimated_coverage(read) < C:
        accept(read)
    else:
        discard(read)
```




Redundant reads
(not needed for assembly)

DIGITAL NORMALIZATION




DIGITAL NORMALIZATION

```
for read in dataset:  
    if estimated_coverage(read) < C:  
        accept(read)  
    else:  
        discard(read)
```



DIGITAL NORMALIZATION

No error



3mer freq.

CAT=32

ATG=34

TGC=36

GCA=35

CAT=33

ATT=34

TTG=40

CATGCATTG

CAT

ATG

TGC

GCA


CAT

ATT

TTG

DIGITAL NORMALIZATION

1error



3mer freq.

CAT=32

ATG=34

TGA=1

GAA=1

AAT=1

ATT=34

TTG=40

CATGAATTG

CAT

ATG

TGA

GAA


AAT

ATT

TTG

DIGITAL NORMALIZATION

>1 error



3mer freq.

CAT=32

ATG=34

TGA=1

GAA=1

AAT=1

ATC=1

TCG=1

CATGAATCG

CAT

ATG

TGA

GAA

AAT

ATC

TCG


DIGITAL NORMALIZATION

Median kmer abundance

0 error: 32,33,34,34,35,36,40

1 error: 1,1,1,32,24,24,40

>1 error: 1,1,1,1,1,32,34



DIGITAL NORMALIZATION

Table 1. Digital normalization to C=20 removes many erroneous k-mers from sequencing data sets. Numbers in parentheses indicate number of true k-mers lost at each step, based on reference.

Data set	True 20-mers	20-mers in reads	20-mers at C=20	% reads kept
Simulated genome	399,981	8,162,813	3,052,007 (-2)	19%
Simulated mRNAseq	48,100	2,466,638 (-88)	1,087,916 (-9)	4.1%
<i>E. coli</i> genome	4,542,150	175,627,381 (-152)	90,844,428 (-5)	11%
Yeast mRNAseq	10,631,882	224,847,659 (-683)	10,625,416 (-6,469)	9.3%
Mouse mRNAseq	43,830,642	709,662,624 (-23,196)	43,820,319 (-13,400)	26.4%

DIGITAL NORMALIZATION

Table 4. Single-pass digital normalization to C=20 reduces computational requirements for transcriptome assembly.

Data set	N reads pre/post	Assembly time pre/post	Assembly memory pre/post
Yeast (Oases)	100m / 9.3m	181 min / 12 min (15.1x)	45.2gb / 8.9gb (5.1x)
Yeast (Trinity)	100m / 9.3m	887 min / 145 min (6.1x)	31.8gb / 10.4gb (3.1x)
Mouse (Oases)	100m / 26.4m	761 min / 73 min (10.4x)	116.0gb / 34.6gb (3.4x)
Mouse (Trinity)	100m / 26.4m	2297 min / 634 min (3.6x)	42.1gb / 36.4gb (1.2x)

Genome Assembly

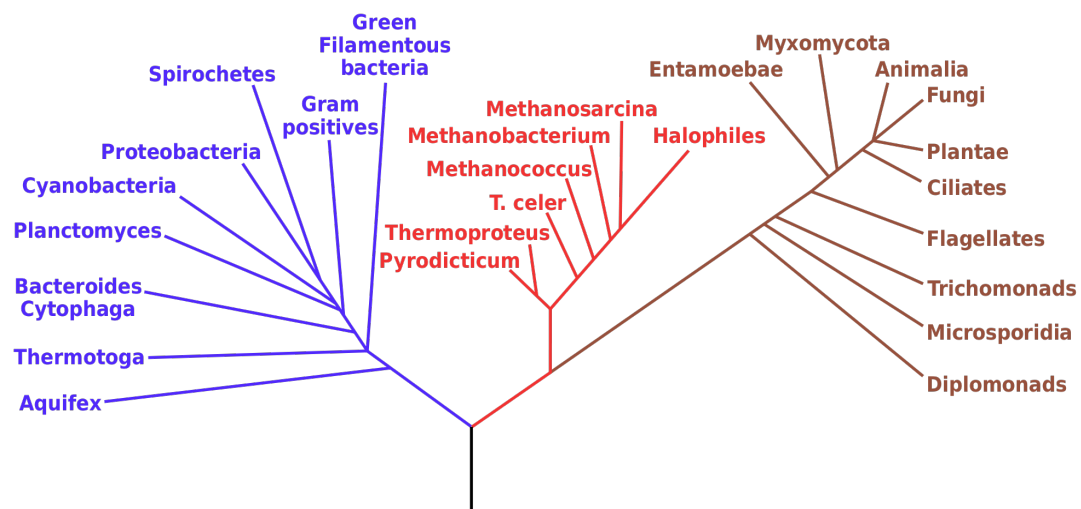
WHY DO YOU WANT TO ASSEMBLE A GENOME?

WHAT DO YOU NEED TO ASSEMBLE A GENOME?

ASSEMBLE A GENOME? GENERAL STRATEGIES


Genome size	Unlimited \$\$	Typical
>10Mb		
10Mb - 100Mb		
> 100 Mb		

GENOME SIZES




ASSEMBLY

Assume sequencing produces such a large # fragments that almost all genome positions are *covered* by many fragments...



ASSEMBLY

...but we don't know what came from where



ASSEMBLY

Key term: *coverage*. Usually it's short for *average coverage*: the average number of reads covering a position in the genome.

CTAGGCCCTCAATTTT
CTTAGGCCCTCAATTTT
GGCTCTAGGCCCTCATTTTT
CTCGGCTCTAGCCCCTCATTTT
TATCTCGACTCTAGGCCCTCA
TATCTCGACTCTAGGCC
TCTATATCTCGGCTCTAGG
GGCGTCTATATCTCG
GGCGTCGATATCT
GGCGTCTATATCT
GGCGTCTATATCTCGGCTCTAGGCCCTCATTTTT

177 nucleotides

35 nucleotides

$$\text{Average coverage} = 177 / 35 \approx 7x$$

OTHER ASSEMBLY TERMS

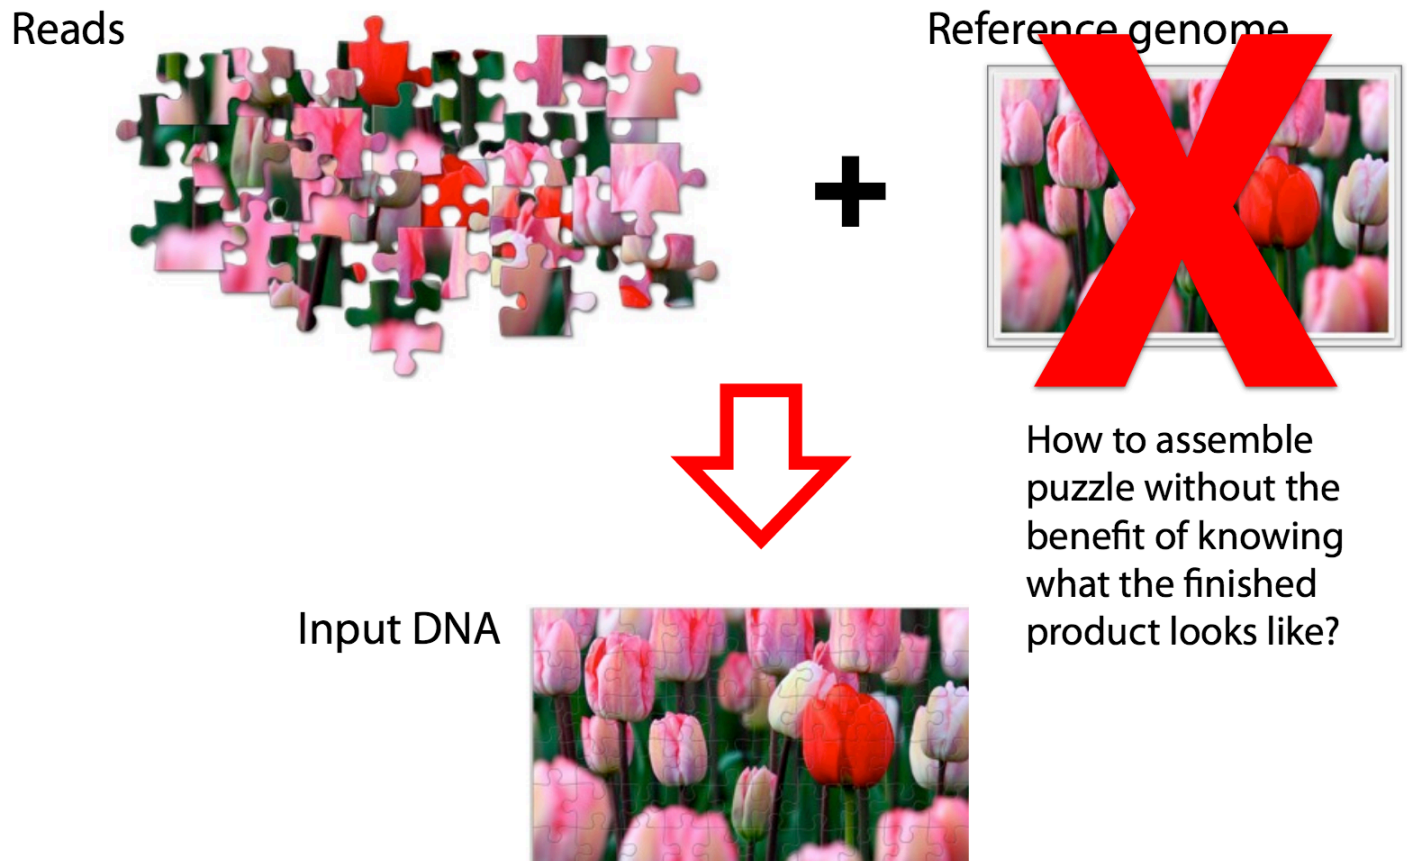
Unitig

Contig

scaffold


ASSEMBLY

- Complicated by:



ASSEMBLY

- Complicated by:



ASSEMBLY

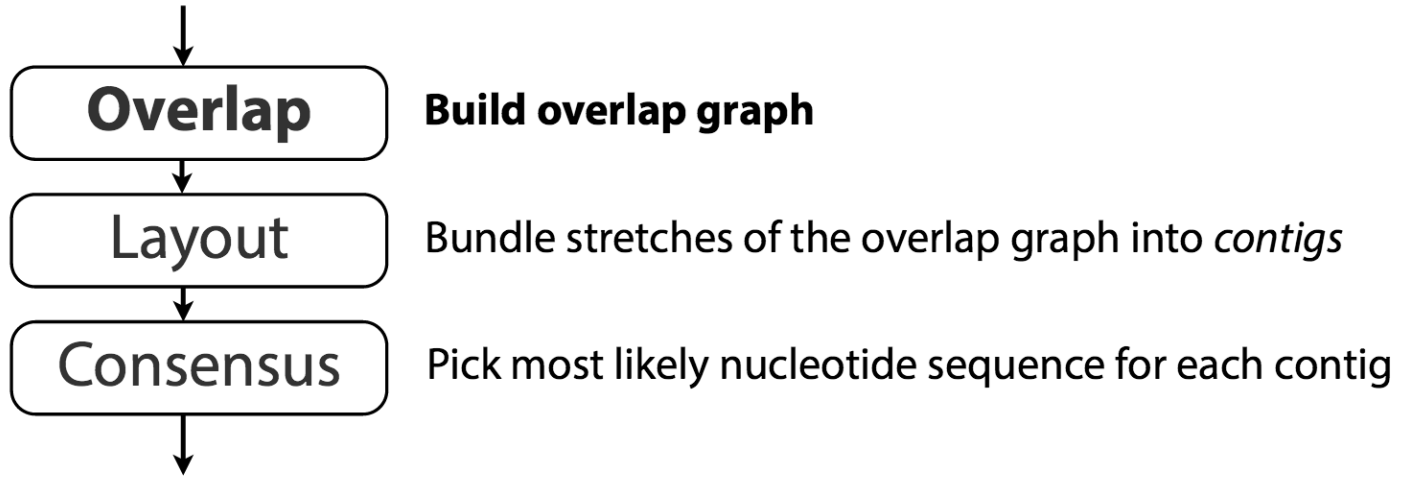
- Work flow:

ASSEMBLY

- 3 assembly strategies:

ASSEMBLY

- OLC Assembly



ASSEMBLY

- OLC Assembly: Characteristics


ASSEMBLY

Directed graph $G(V, E)$ consists of set of *vertices*, V and set of *directed edges*, E

Directed edge is an *ordered pair* of vertices.
First is the *source*, second is the *sink*.

Vertex is drawn as a circle

Edge is drawn as a line with an arrow connecting two circles



Vertex also called *node* or *point*

Edge also called *arc* or *line*


Directed graph also called *digraph*

$$V = \{a, b, c, d\}$$

$$E = \{(a, b), (a, c), (c, b)\}$$


Source Sink

ASSEMBLY



Build overlap graph

to_every_thing_turn_turn_turn_there_is_a_season
L=4, k=7



ASSEMBLY

Overlap

Build overlap graph

Vertices (reads): { *a*: CTCTAGGCC, *b*: GCCCTCAAT, *c*: CAATTTT }

Edges (overlaps): { (*a*, *b*), (*b*, *c*) }

a: CTCTAGGCC

3

b: GCCCTCAAT

4

c: CAATTTT

CTCTAGGCC

|||

GCCCTCAAT

GCCCTCAAT

||||

CAATTTT